

```
In [1]: import os
import numpy as np
import pandas as pd
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: df=pd.read_csv('C:/Users/User/Downloads/heart.csv')
```

```
In [4]: df.head()
```

```
Out[4]:
```


	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
In [5]: plt.figure(figsize=(12,8))
sns.heatmap(df.corr(),square=False,cmap='flare',annot=True)
```

Out[5]: <Axes: >



```
In [7]: #creating features (x) and targets(y) variables
x=df.iloc[:, :-1]
y=df.iloc[:, -1]
```

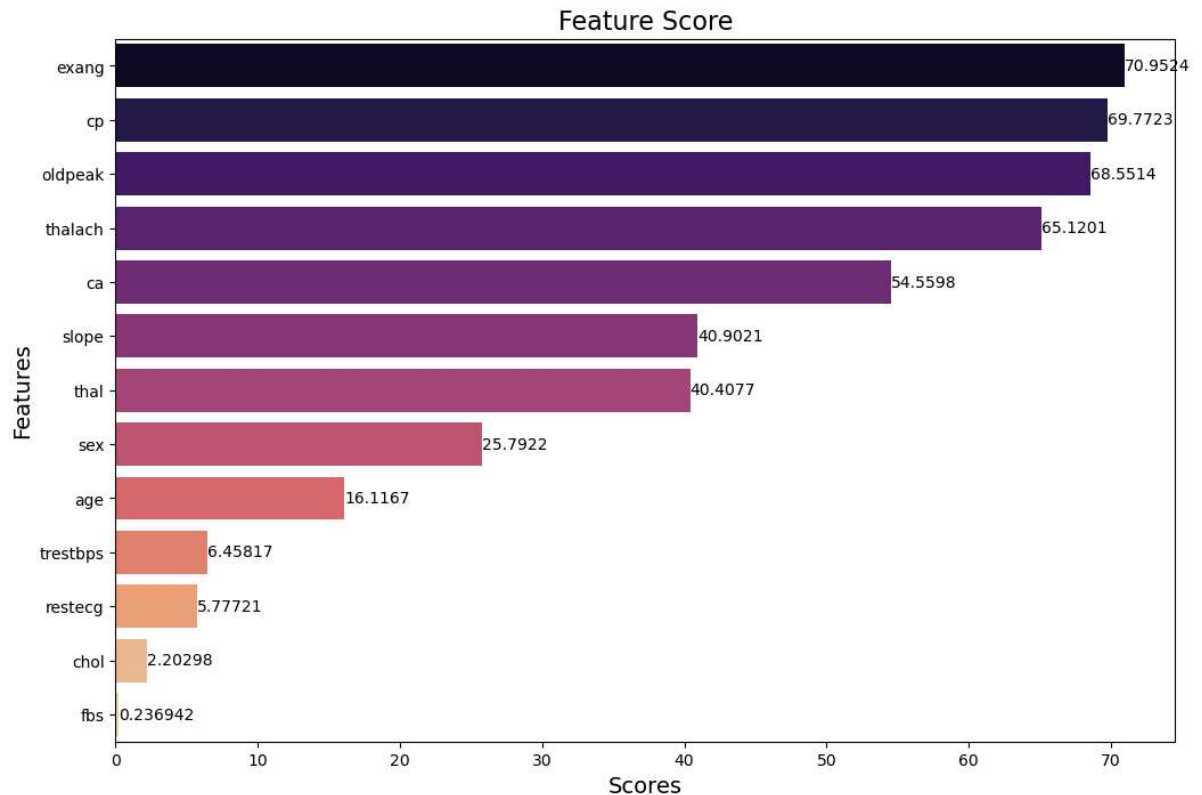
```
In [15]: #Finding the best features for the model
from sklearn.feature_selection import SelectKBest, f_classif
selector=SelectKBest(f_classif,k=13)
x_selected=selector.fit_transform(x,y)
```

```
In [18]: #List of all column ans score inx variable
selected_features=x.columns[selector.get_support()]
feature_scores=selector.scores_[selector.get_support()]
```

```
In [19]: #Creating dataframe for feature and scores
feature_score_df=pd.DataFrame({'Features': selected_features, 'Scores': feature_scores})
```

```
In [21]: #sort the dataframe in descending order
feature_score_df=feature_score_df.sort_values(by='Scores', ascending=False)
```

```
In [25]: #plotting a barplot for better understanding
plt.figure(figsize=(12,8))
ax=sns.barplot(x=feature_score_df['Scores'],y=feature_score_df['Features'], palette='magma')
plt.title('Feature Score', fontsize=16)
plt.xlabel('Scores',fontsize=14)
plt.ylabel('Features',fontsize=14)
for lab in ax.containers:
    ax.bar_label(lab)
```



```
In [26]: #dropping the low scored values
x=x.drop(['fbs','chol','restecg','trestbps'],axis=1)
```

```
In [31]: from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
x=scaler.fit_transform(x)
```

In [29]: `x.head`

Out[29]: <bound method NDFrame.head of

	exang	oldpeak	slope \	age	sex	cp	thalach
0	0.952197	0.681005	1.973123	0.015443	-0.696631	1.087338	-2.274579
1	-1.915313	0.681005	1.002577	1.633471	-0.696631	2.122573	-2.274579
2	-1.474158	-1.468418	0.032031	0.977514	-0.696631	0.310912	0.976352
3	0.180175	0.681005	0.032031	1.239897	-0.696631	-0.206705	0.976352
4	0.290464	-1.468418	-0.938515	0.583939	1.435481	-0.379244	0.976352
..
298	0.290464	-1.468418	-0.938515	-1.165281	1.435481	-0.724323	-0.649113
299	-1.033002	0.681005	1.973123	-0.771706	-0.696631	0.138373	-0.649113
300	1.503641	0.681005	-0.938515	-0.378132	-0.696631	2.036303	-0.649113
301	0.290464	0.681005	-0.938515	-1.515125	1.435481	0.138373	-0.649113
302	0.290464	-1.468418	0.032031	1.064975	-0.696631	-0.896862	-0.649113

	ca	thal
0	-0.714429	-2.148873
1	-0.714429	-0.512922
2	-0.714429	-0.512922
3	-0.714429	-0.512922
4	-0.714429	-0.512922
..
298	-0.714429	1.123029
299	-0.714429	1.123029
300	1.244593	1.123029
301	0.265082	1.123029
302	0.265082	-0.512922

[303 rows x 9 columns]>

In [32]: `from sklearn.model_selection import train_test_split`
`x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=4`

In [37]: `from sklearn.neighbors import KNeighborsClassifier`
`knn=KNeighborsClassifier(n_neighbors=11)`
`knn.fit(x_train,y_train)`

Out[37]: KNeighborsClassifier(n_neighbors=11)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [38]: `y_predict=knn.predict(x_test)`

In [39]: `from sklearn.metrics import accuracy_score`
`print(accuracy_score(y_test,y_predict))`

0.9016393442622951


```
In [40]: from sklearn.metrics import classification_report  
print(classification_report(y_test,y_predict))
```

	precision	recall	f1-score	support
0	0.90	0.90	0.90	29
1	0.91	0.91	0.91	32
accuracy			0.90	61
macro avg	0.90	0.90	0.90	61
weighted avg	0.90	0.90	0.90	61

```
In [42]: from sklearn.metrics import confusion_matrix  
cm=confusion_matrix(y_test,y_predict)  
print(cm)
```

```
[[26  3]  
 [ 3 29]]
```

```
In [47]: plt.figure(figsize=(9,9))
plt.title('Confusion matrix')
sns.heatmap(cm,annot=True,fmt='g',linewidth=0.5,cmap='Blues_r',)
plt.xlabel('Predicted label',fontsize=12)
plt.ylabel('Actual label',fontsize=12)
plt.show()
```



```
In [ ]:
```